


Benefits of Concurrency in Java & Android: Performance



Douglas C. Schmidt

d.schmidt@vanderbilt.edu

www.dre.vanderbilt.edu/~schmidt

Institute for Software
Integrated Systems
Vanderbilt University
Nashville, Tennessee, USA



Learning Objectives in this Part of the Module

- Recognize how concurrency can improve performance in Java & Android



Using Concurrency to Increase Performance

Using Concurrency to Increase Performance

- “Performance” is a characterization of the amount of useful work that can be accomplished



See [en.wikipedia.org/wiki/
Computer_performance](https://en.wikipedia.org/wiki/Computer_performance)

Using Concurrency to Increase Performance

- “Performance” is a characterization of the amount of useful work that can be accomplished, e.g.
- Decreasing response time for handling requests



See en.wikipedia.org/wiki/Computer_performance#Response_time

Using Concurrency to Increase Performance

- “Performance” is a characterization of the amount of useful work that can be accomplished, e.g.
- Decreasing response time for handling requests

- *Service time* – How long it takes to do requested work
- *Wait time* – How long the request has to wait for before it gets to run
- *Transmission time* – How long it takes to move request to computer doing the work & the response back to requestor



Using Concurrency to Increase Performance

- “Performance” is a characterization of the amount of useful work that can be accomplished, e.g.
- Decreasing response time for handling requests

- *Service time* – How long it takes to do requested work
- *Wait time* – How long the request has to wait for before it gets to run
- *Transmission time* – How long it takes to move request to computer doing the work & the response back to requestor



Using Concurrency to Increase Performance

- “Performance” is a characterization of the amount of useful work that can be accomplished, e.g.
- Decreasing response time for handling requests

- *Service time* – How long it takes to do requested work
- *Wait time* – How long the request has to wait for before it gets to run
- *Transmission time* – How long it takes to move request to computer doing the work & the response back to requestor



Using Concurrency to Increase Performance

- “Performance” is a characterization of the amount of useful work that can be accomplished, e.g.
 - Decreasing response time for handling requests
 - Increasing the amount of work that can be performed within a given time



See en.wikipedia.org/wiki/Computer_performance#Throughput

Using Concurrency to Increase Performance

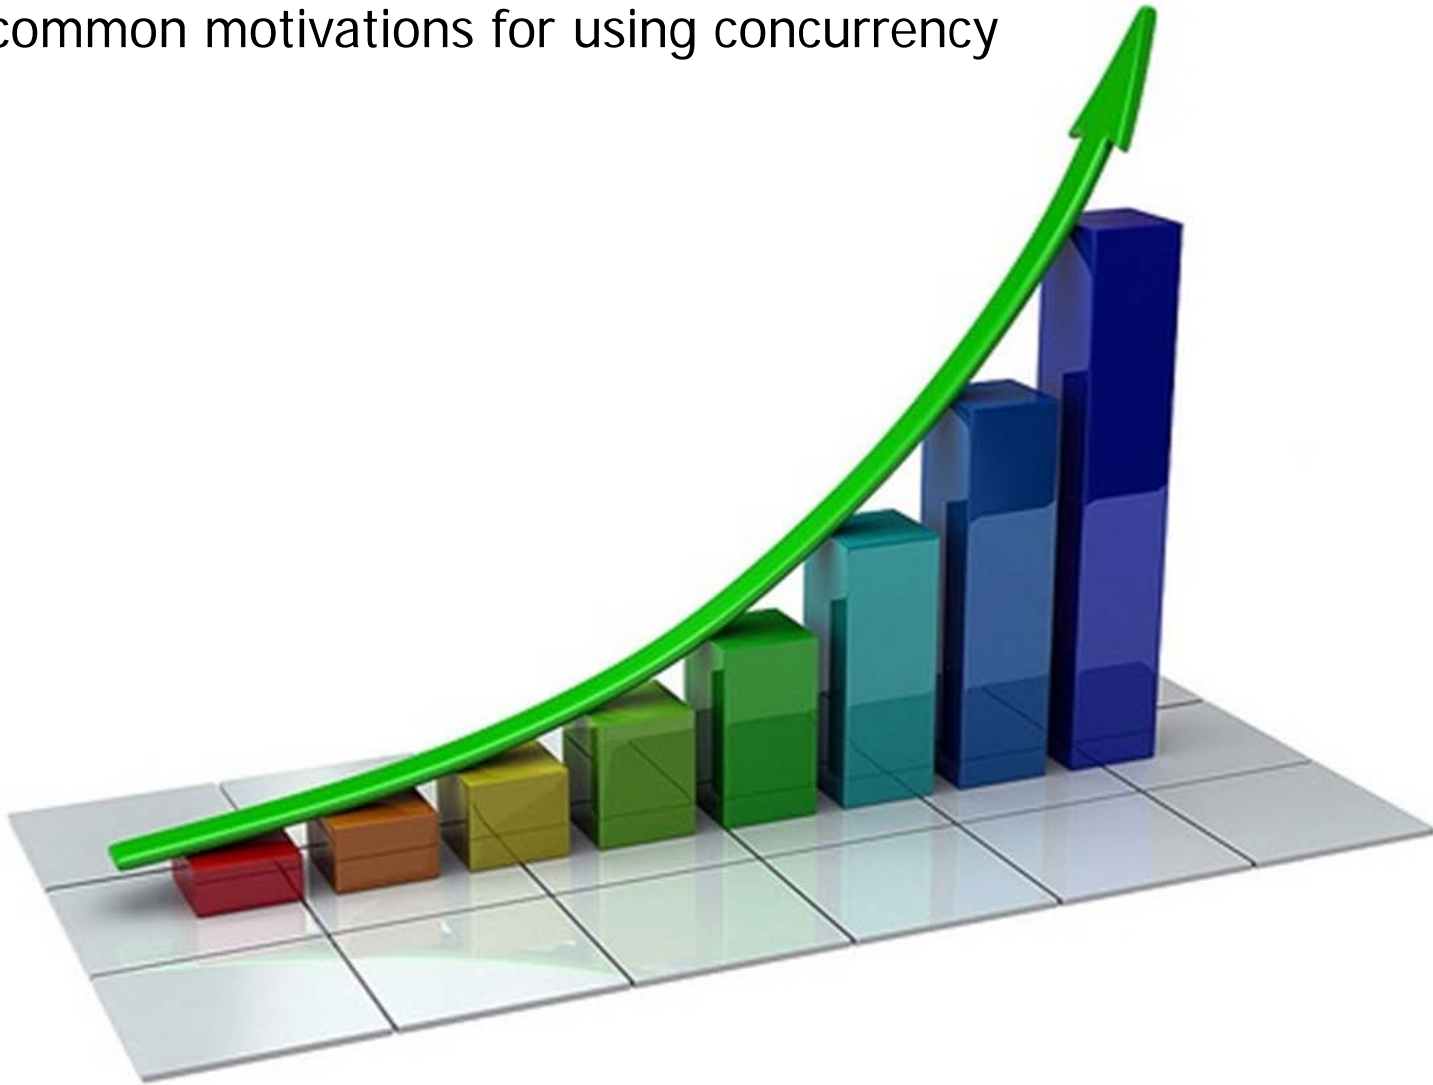
- “Performance” is a characterization of the amount of useful work that can be accomplished, e.g.
 - Decreasing response time for handling requests
 - Increasing the amount of work that can be performed within a given time



See en.wikipedia.org/wiki/Up_to_eleven
for more on maximizing performance


Using Concurrency to Increase Performance

- Decreasing response time & increasing the amount of work performed within a given time are common motivations for using concurrency



Using Concurrency to Increase Performance

- Decreasing response time & increasing the amount of work performed within a given time are common motivations for using concurrency



Optimizing these properties requires an understanding of patterns & tradeoffs amongst constraints & quality attributes

Using Concurrency to Increase Performance

- Performance can be accelerated via parallel processing




See [en.wikipedia.org/
wiki/Parallel_computing](http://en.wikipedia.org/wiki/Parallel_computing)

Using Concurrency to Increase Performance

- Performance can be accelerated via parallel processing, e.g.
- Performing computations simultaneously




→3 →3 →3 →3 →3



Using Concurrency to Increase Performance


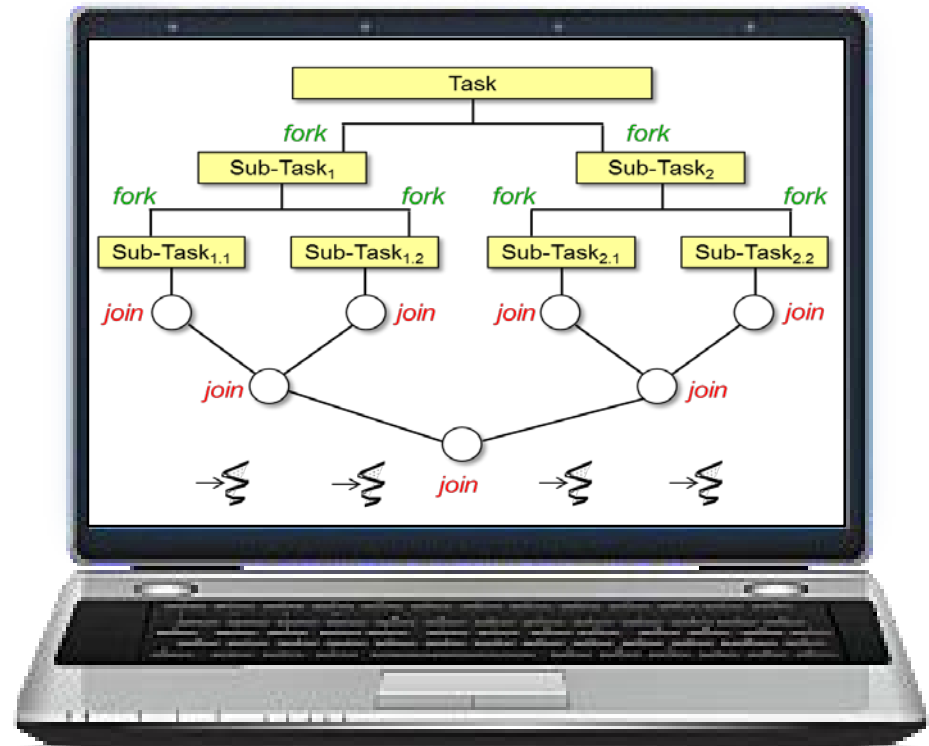
- Performance can be accelerated via parallel processing, e.g.
 - Performing computations simultaneously
 - Particularly for computations with no inter-dependencies



See en.wikipedia.org/wiki/Embarrassingly_parallel

Using Concurrency to Increase Performance


- Performance can be accelerated via parallel processing, e.g.
 - Performing computations simultaneously
- Dividing a large problem into multiple smaller problems that can be processed in parallel



See en.wikipedia.org/wiki/Divide_and_conquer_algorithm

Using Concurrency to Increase Performance


- Performance can be accelerated via parallel processing, e.g.
 - Performing computations simultaneously
 - Dividing a large problem into multiple smaller problems that can be processed in parallel, e.g.
 - Web searches



See en.wikipedia.org/wiki/MapReduce

Using Concurrency to Increase Performance


- Performance can be accelerated via parallel processing, e.g.
 - Performing computations simultaneously
 - Dividing a large problem into multiple smaller problems that can be processed in parallel, e.g.
 - Web searches
 - Image rendering



See en.wikipedia.org/wiki/Parallel_rendering

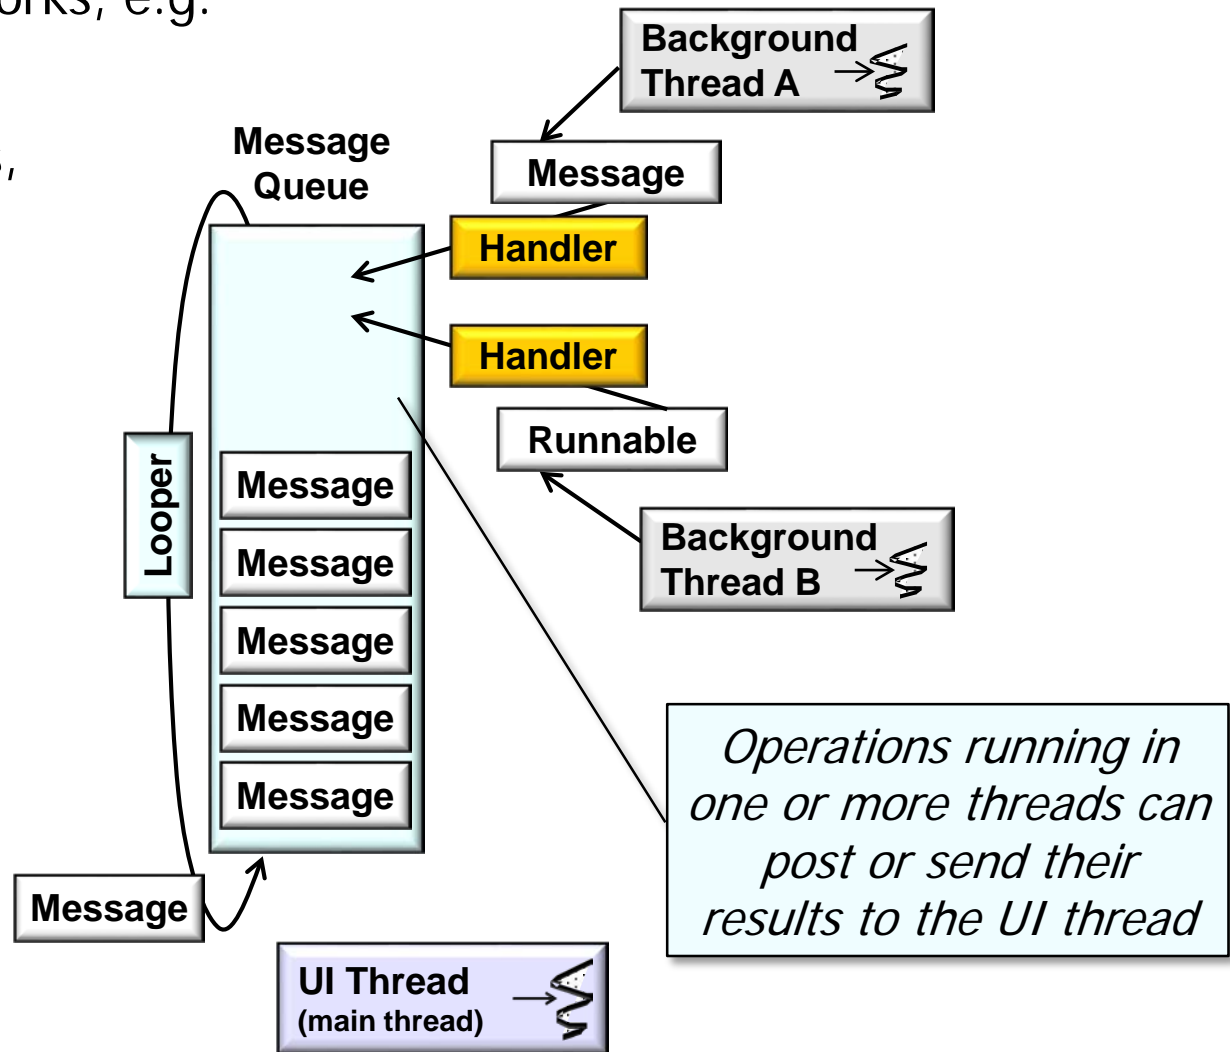
Using Concurrency to Increase Performance

- Android enables parallelism by overlapping computation & communication via two concurrency frameworks



Using Concurrency to Increase Performance

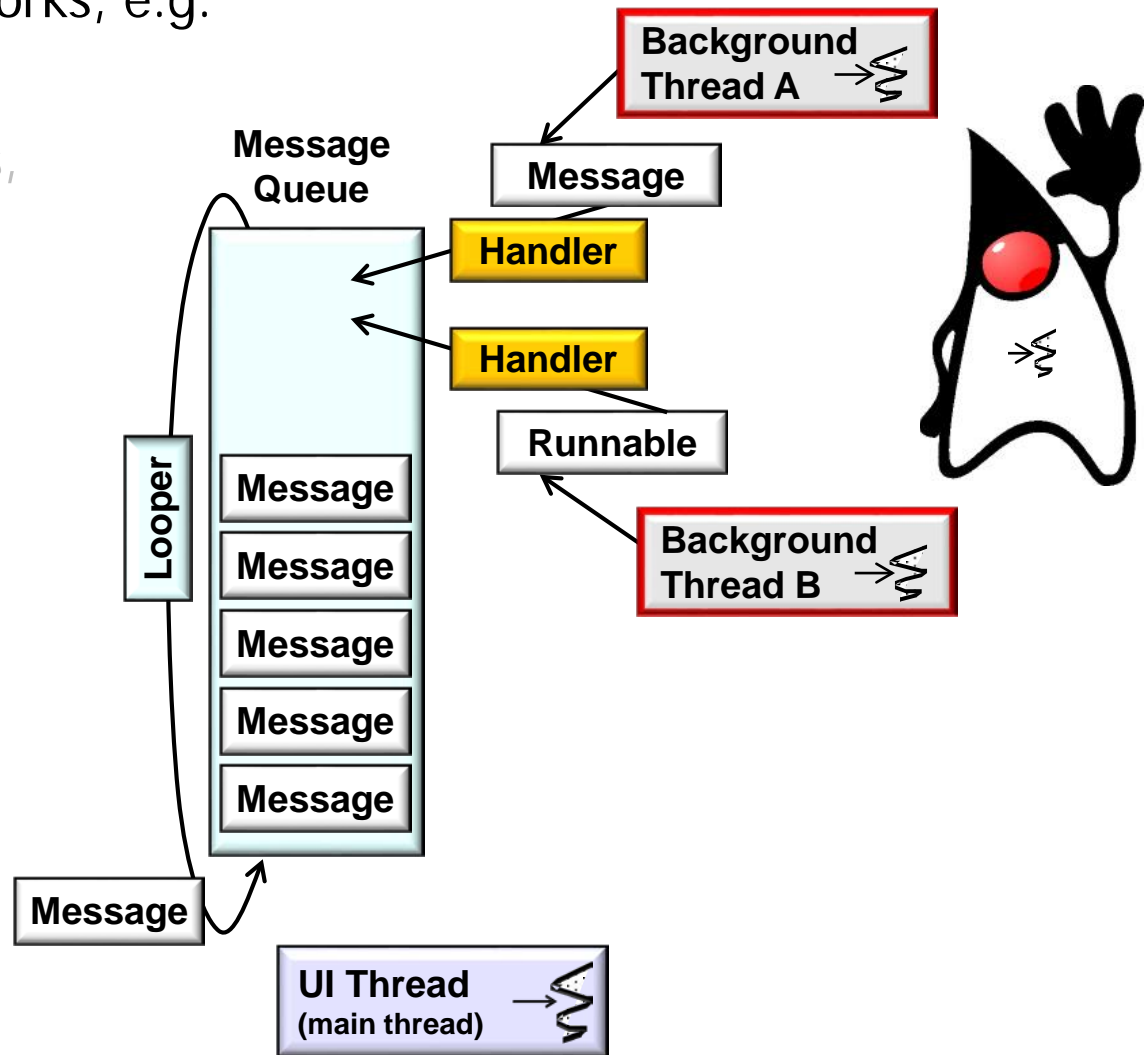
- Android enables parallelism by overlapping computation & communication via two concurrency frameworks, e.g.
 - HaMeR framework
 - “Handlers, Messages, & Runnables”



See code.tutsplus.com/tutorials/concurrency-on-android-using-hamer-framework--cms-27129

Using Concurrency to Increase Performance

- Android enables parallelism by overlapping computation & communication via two concurrency frameworks, e.g.
 - HaMeR framework
 - “Handlers, Messages, & Runnables”
 - Works with Java threads



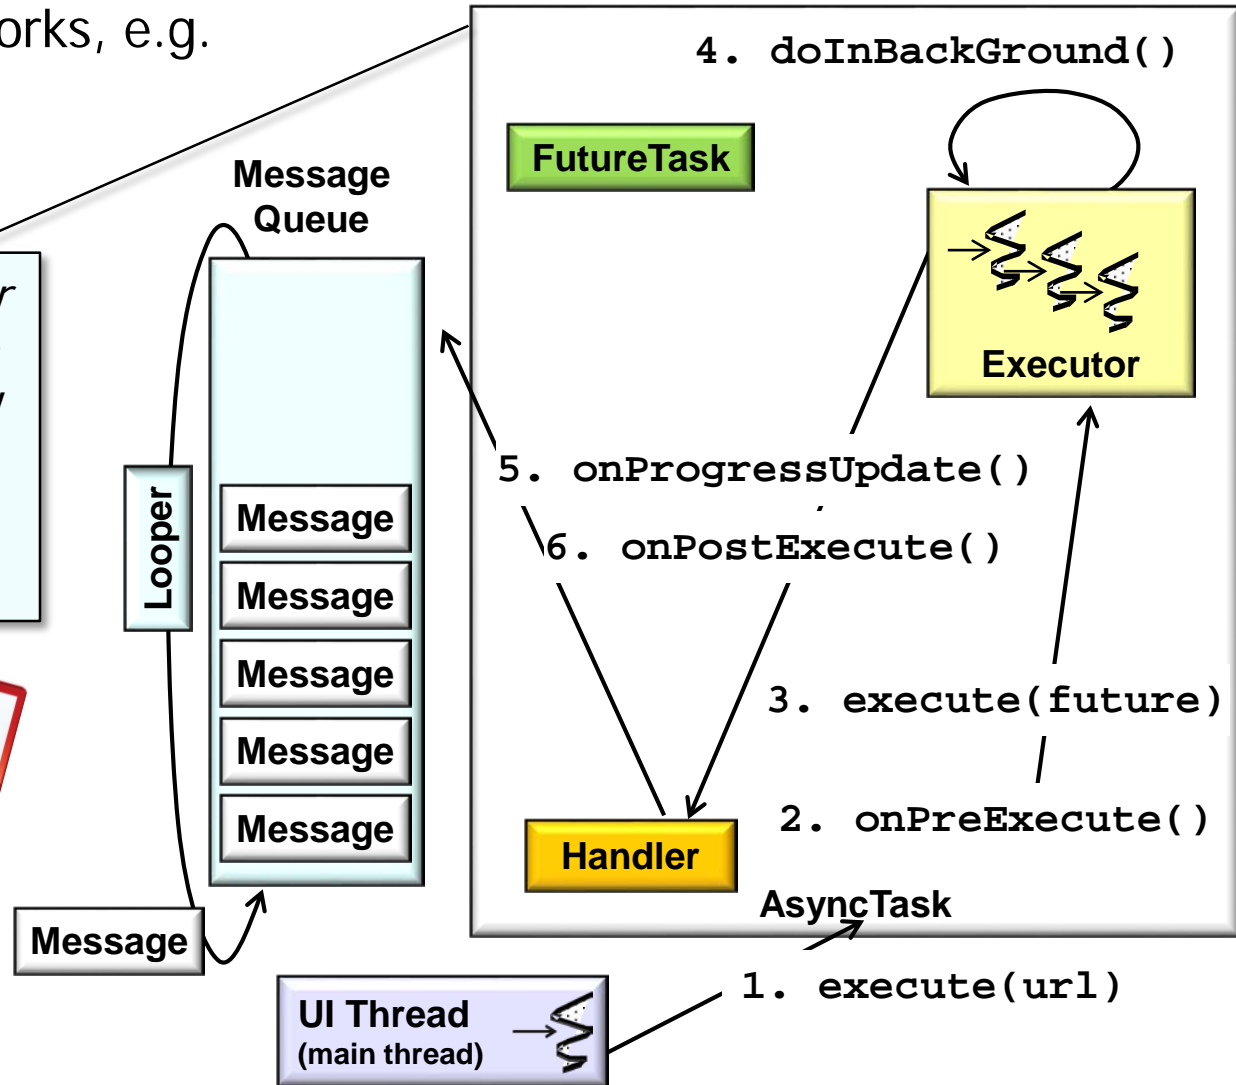
See docs.oracle.com/javase/tutorial/essential/concurrency/threads.html

Using Concurrency to Increase Performance

- Android enables parallelism by overlapping computation & communication via two concurrency frameworks, e.g.

- HaMeR framework
- AsyncTask framework

Operations run in one or more threads & publish results to the UI thread without using threads, handlers, messages, and/or runnables

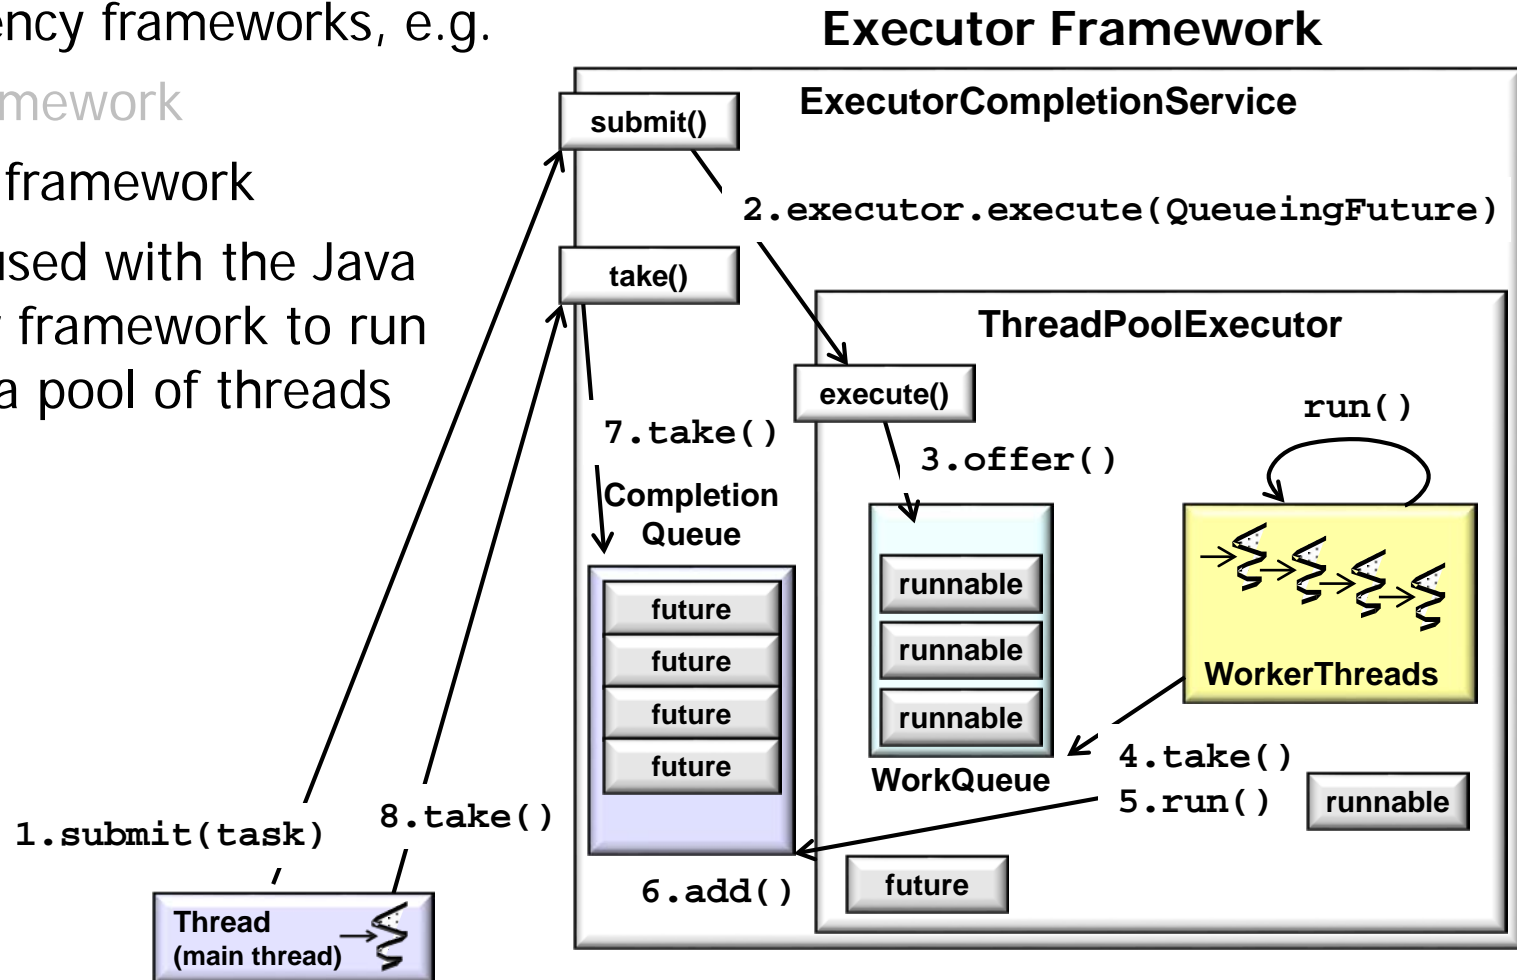


See developer.android.com/reference/android/os/AsyncTask.html

Using Concurrency to Increase Performance

- Android enables parallelism by overlapping computation & communication via two concurrency frameworks, e.g.

- HaMeR framework
- AsyncTask framework
 - Can be used with the Java Executor framework to run tasks in a pool of threads



See docs.oracle.com/javase/tutorial/essential/concurrency/executors.html

End Benefits of Concurrency in Java & Android (Part 1)

Benefits of Concurrency in Java & Android: Responsiveness



Douglas C. Schmidt

d.schmidt@vanderbilt.edu

www.dre.vanderbilt.edu/~schmidt

Institute for Software
Integrated Systems
Vanderbilt University
Nashville, Tennessee, USA



Learning Objectives in this Part of the Module

- Recognize how concurrency can improve performance in Java & Android
- Recognize how concurrency can improve responsiveness in Java & Android



Impediments to Leveraging Hardware Parallelism

Impediments to Leveraging Hardware Parallelism

- Despite all the benefits of concurrency to improve performance, there are limits in practice to leveraging hardware parallelism





Impediments to Leveraging Hardware Parallelism

- Not every computing platform supports the latest hardware advances



Impediments to Leveraging Hardware Parallelism

- Not every computing platform supports the latest hardware advances
 - e.g., older computing devices just have a single core, which limits available parallelism



Impediments to Leveraging Hardware Parallelism

- It's also hard to fully leverage parallelism due to various impediments, e.g.



Impediments to Leveraging Hardware Parallelism

- It's also hard to fully leverage parallelism due to various impediments, e.g.
 - Legacy code that's not thread safe

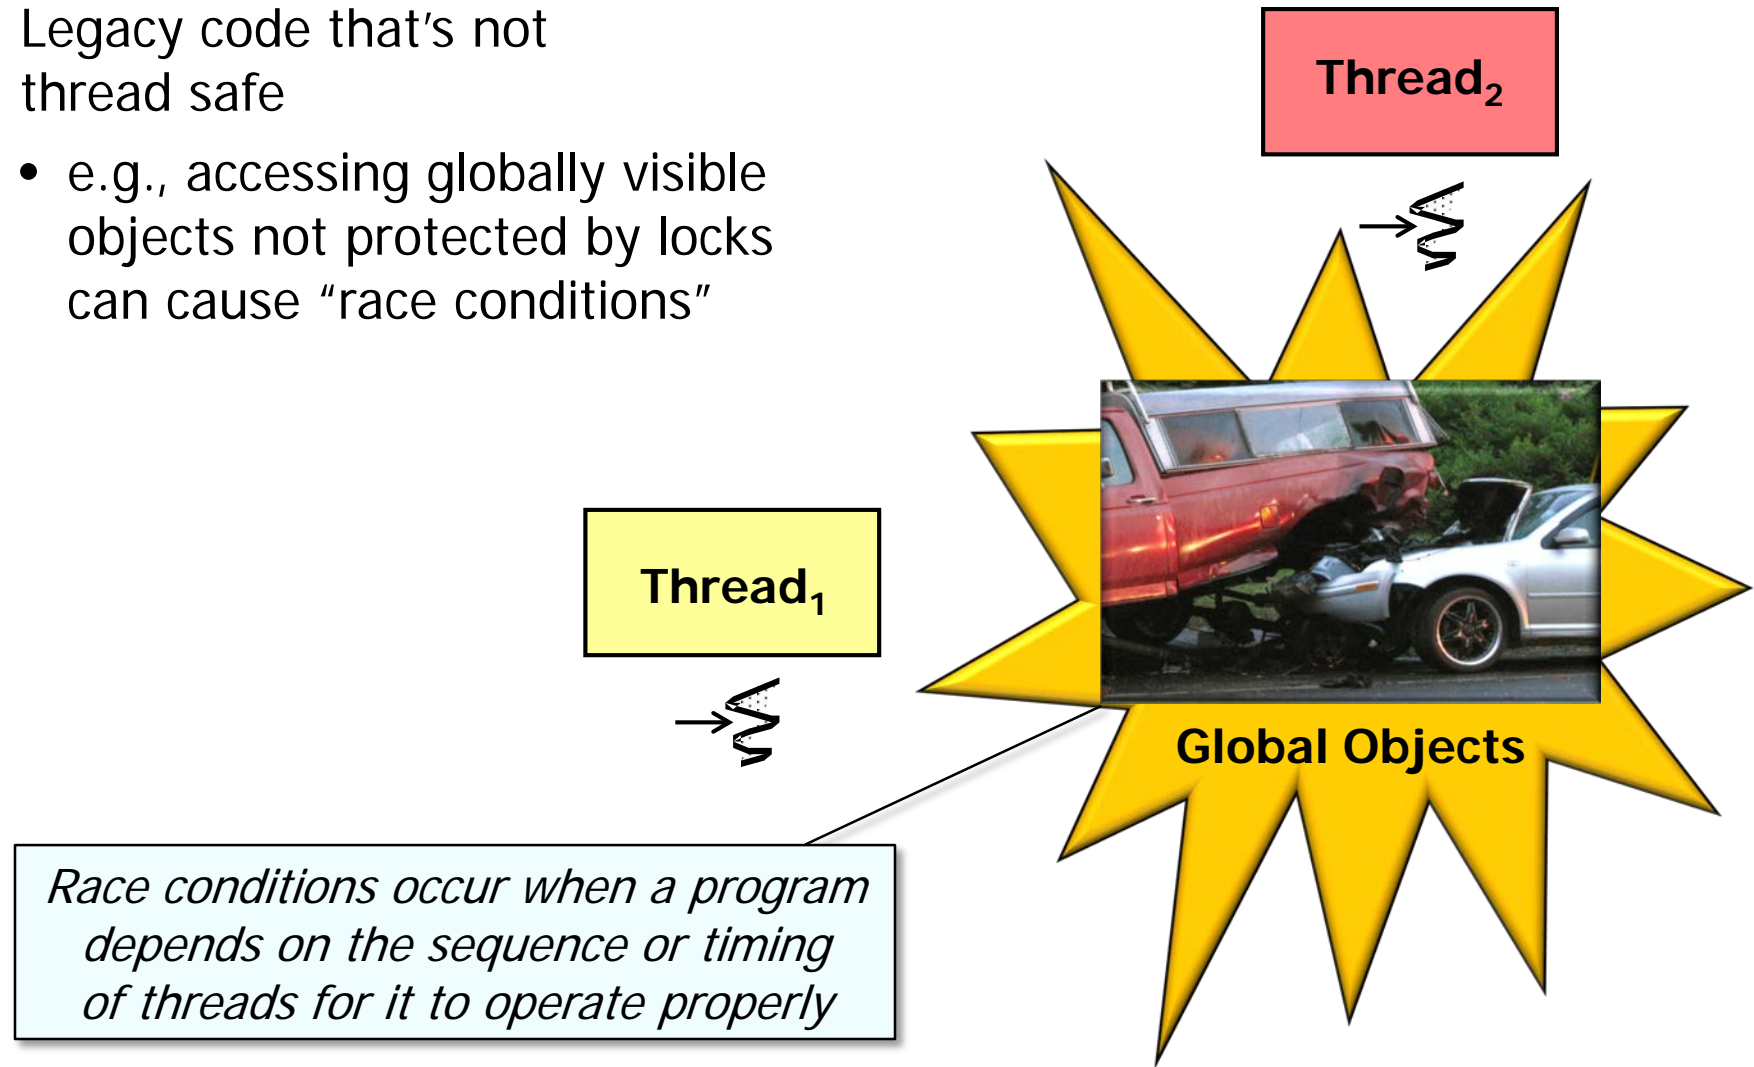


See en.wikipedia.org/wiki/Legacy_system

See www.wesleysteiner.com/professional/MakingLegacyCodeSafe.html

Impediments to Leveraging Hardware Parallelism


- It's also hard to fully leverage parallelism due to various impediments, e.g.
 - Legacy code that's not thread safe
 - e.g., accessing globally visible objects not protected by locks can cause "race conditions"



See en.wikipedia.org/wiki/Race_condition#Software

Impediments to Leveraging Hardware Parallelism

- It's also hard to fully leverage parallelism due to various impediments, e.g.
 - Legacy code that's not thread safe
 - GUI toolkits aren't thread-safe by design



See community.oracle.com/blogs/kgb/2004/10/19/multithreaded-toolkits-failed-dream

Impediments to Leveraging Hardware Parallelism

- It's also hard to fully leverage parallelism due to various impediments, e.g.
 - Legacy code that's not thread safe
 - GUI toolkits aren't thread-safe by design, e.g.
 - Eliminate the need for internal locking



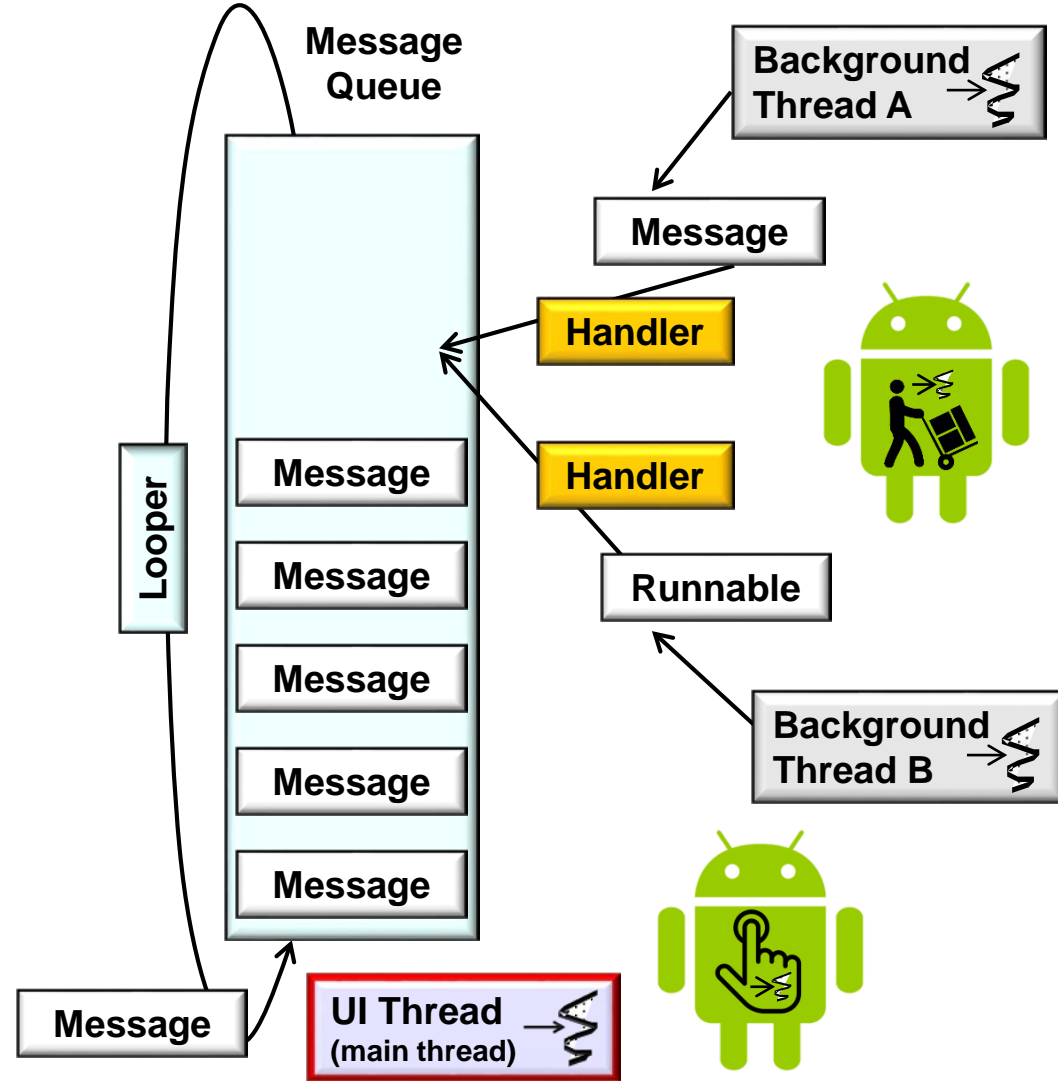
Impediments to Leveraging Hardware Parallelism

- It's also hard to fully leverage parallelism due to various impediments, e.g.
 - Legacy code that's not thread safe
 - GUI toolkits aren't thread-safe by design, e.g.
 - Eliminate the need for internal locking
 - Minimize the need for app developers to understand concurrency



Impediments to Leveraging Hardware Parallelism

- It's also hard to fully leverage parallelism due to various impediments, e.g.
 - Legacy code that's not thread safe
 - GUI toolkits aren't thread-safe by design, e.g.
 - Eliminate the need for internal locking
 - Minimize the need for app developers to understand concurrency
 - Android only allows the UI thread to access GUI components




See developer.android.com/training/multiple-threads/communicate-ui.html

Impediments to Leveraging Hardware Parallelism

- It's also hard to fully leverage parallelism due to various impediments, e.g.
 - Legacy code that's not thread safe
 - GUI toolkits aren't thread-safe by design
 - The impact of Amdahl's Law

"The speedup of a program using multiple processors is limited by the sequential portion of the program that can't run in parallel"



See en.wikipedia.org/wiki/Amdahl's_law

Using Concurrency to Improve Responsiveness

Using Concurrency to Improve Responsiveness

- Concurrency can often be used to improve *perceived* response time



See [en.wikipedia.org/
wiki/Responsiveness](https://en.wikipedia.org/wiki/Responsiveness)

Using Concurrency to Improve Responsiveness

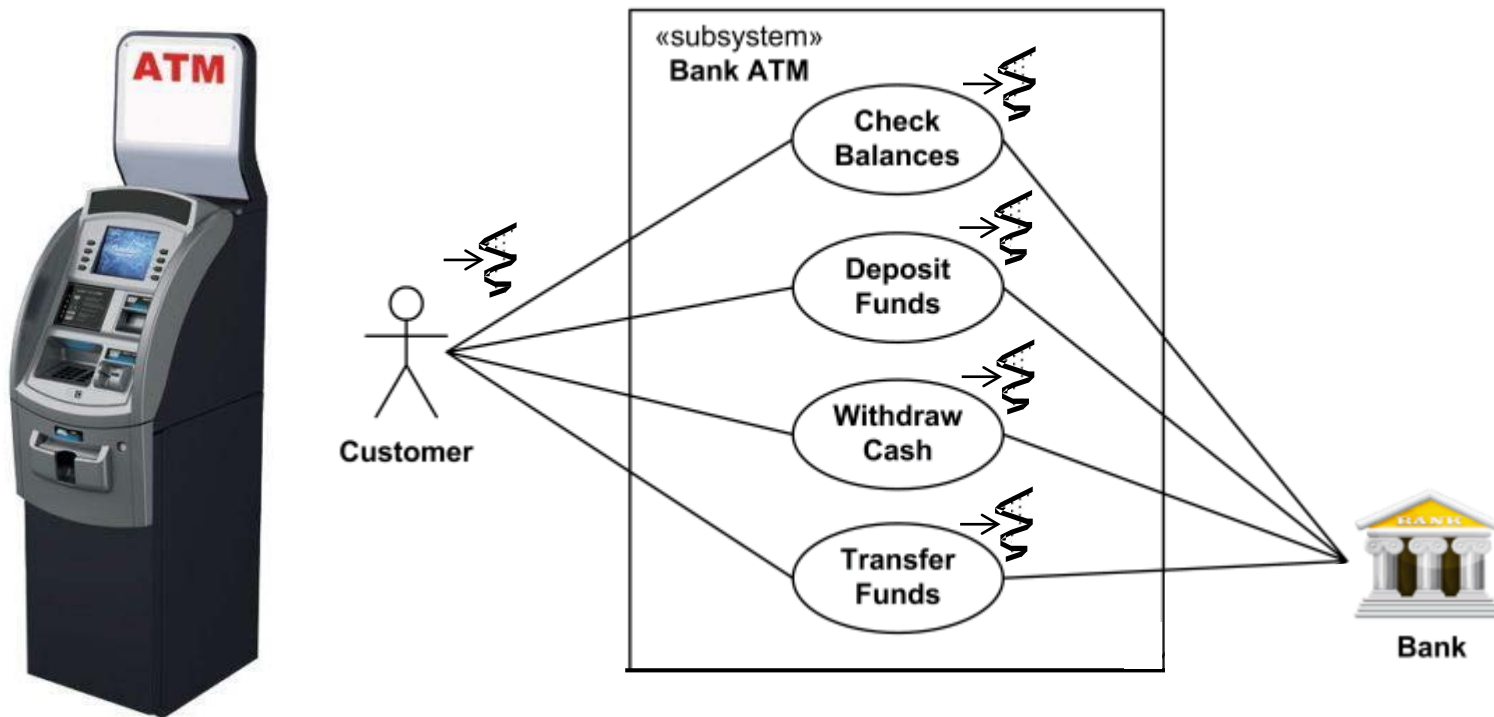
- Concurrency can often be used to improve *perceived* response time, e.g.
 - Don't ignore user input while long-duration computations or communications are occurring



See en.wikipedia.org/wiki/Spinning_pinwheel
& en.wikipedia.org/wiki/Windows_wait_cursor


Using Concurrency to Improve Responsiveness

- Concurrency can often be used to improve *perceived* response time, e.g.
 - Don't ignore user input while long-duration computations or communications are occurring
 - e.g., allow worker threads to perform other processing in the background, while another thread handles user input



Using Concurrency to Improve Responsiveness

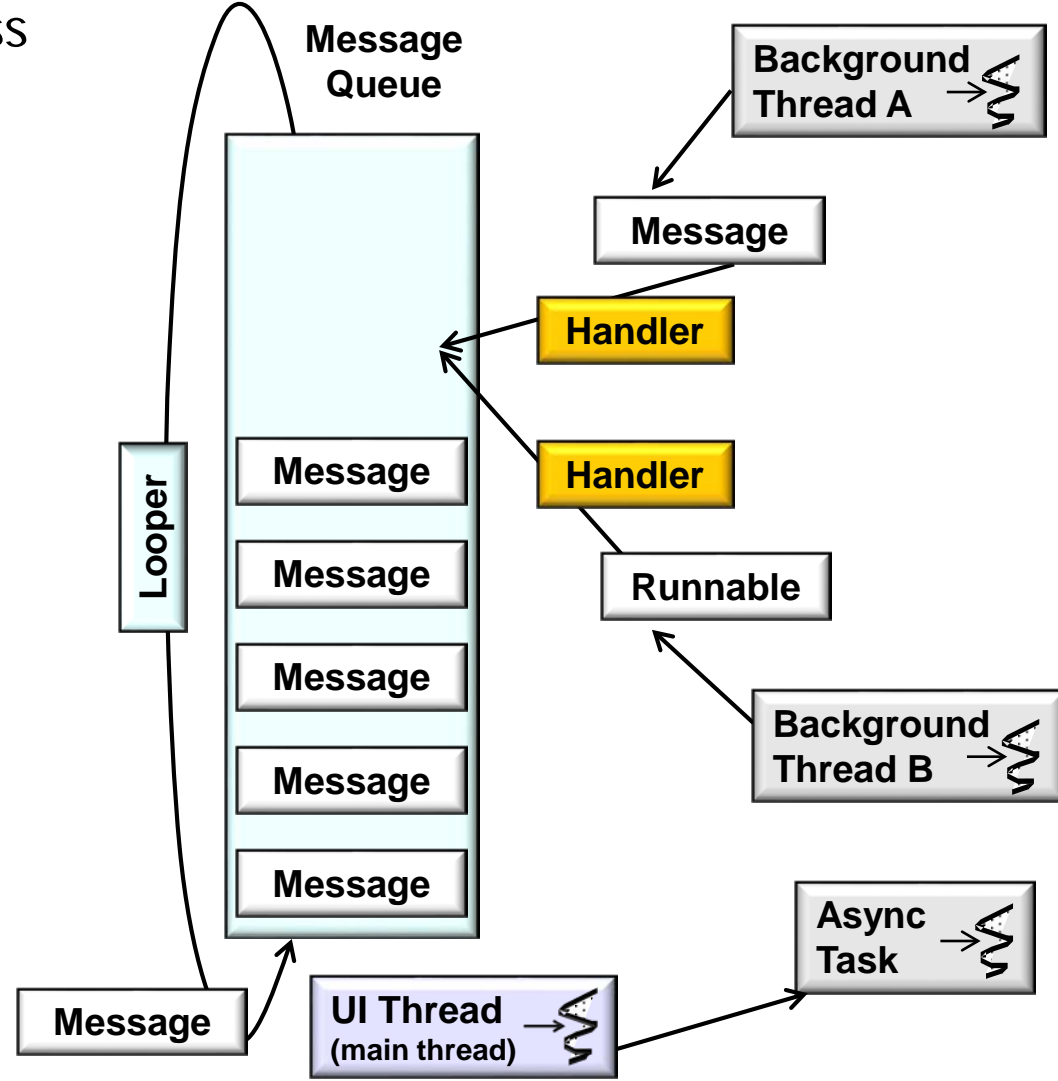
- Concurrency can often be used to improve *perceived* response time, e.g.
 - Don't ignore user input while long-duration computations or communications are occurring
- As long as the software infrastructure supports preemptive multi-threading even single core hardware can be more responsive



See [en.wikipedia.org/wiki/Preemption_\(computing\)](https://en.wikipedia.org/wiki/Preemption_(computing))

Using Concurrency to Improve Responsiveness

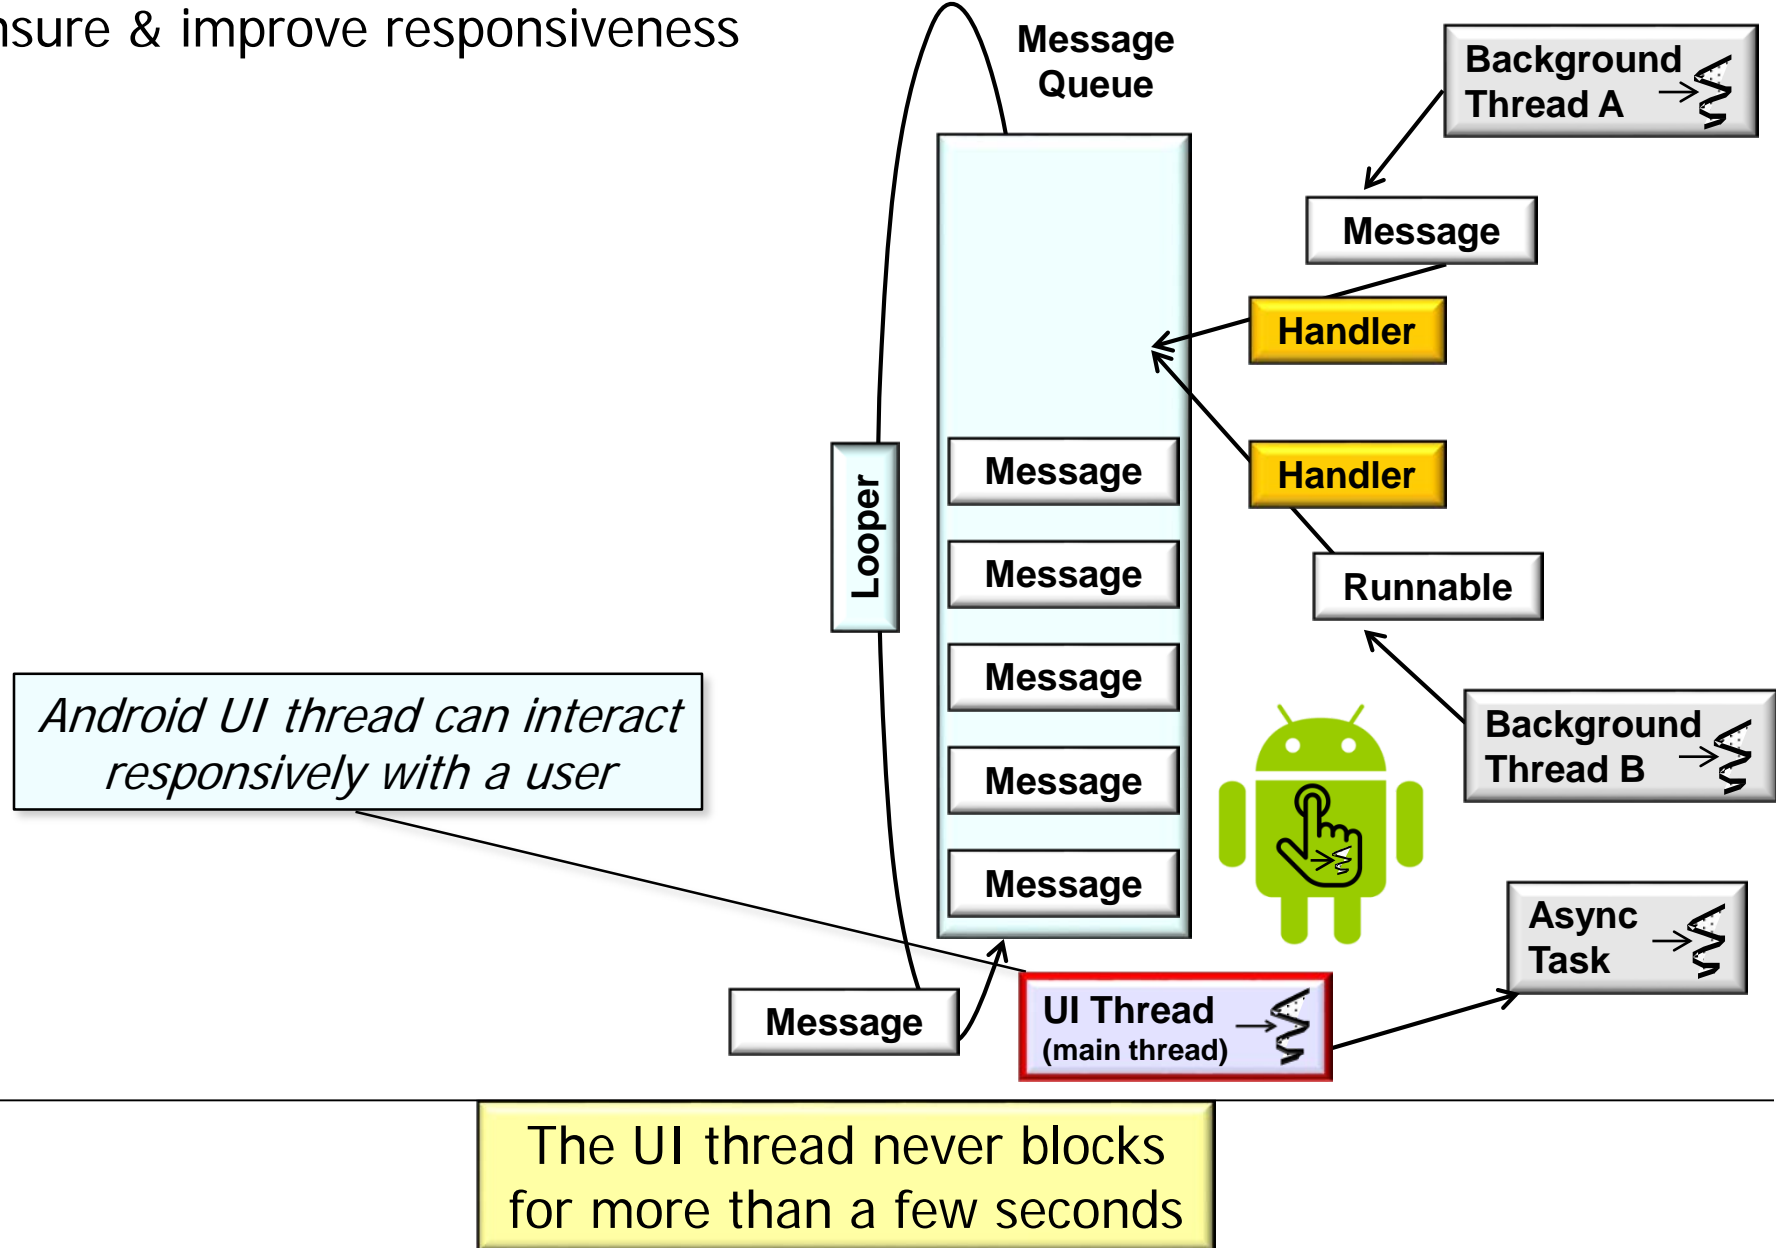
- Android concurrency frameworks define features & idioms that can help ensure & improve responsiveness



See developer.android.com/training/articles/perf-anr.html

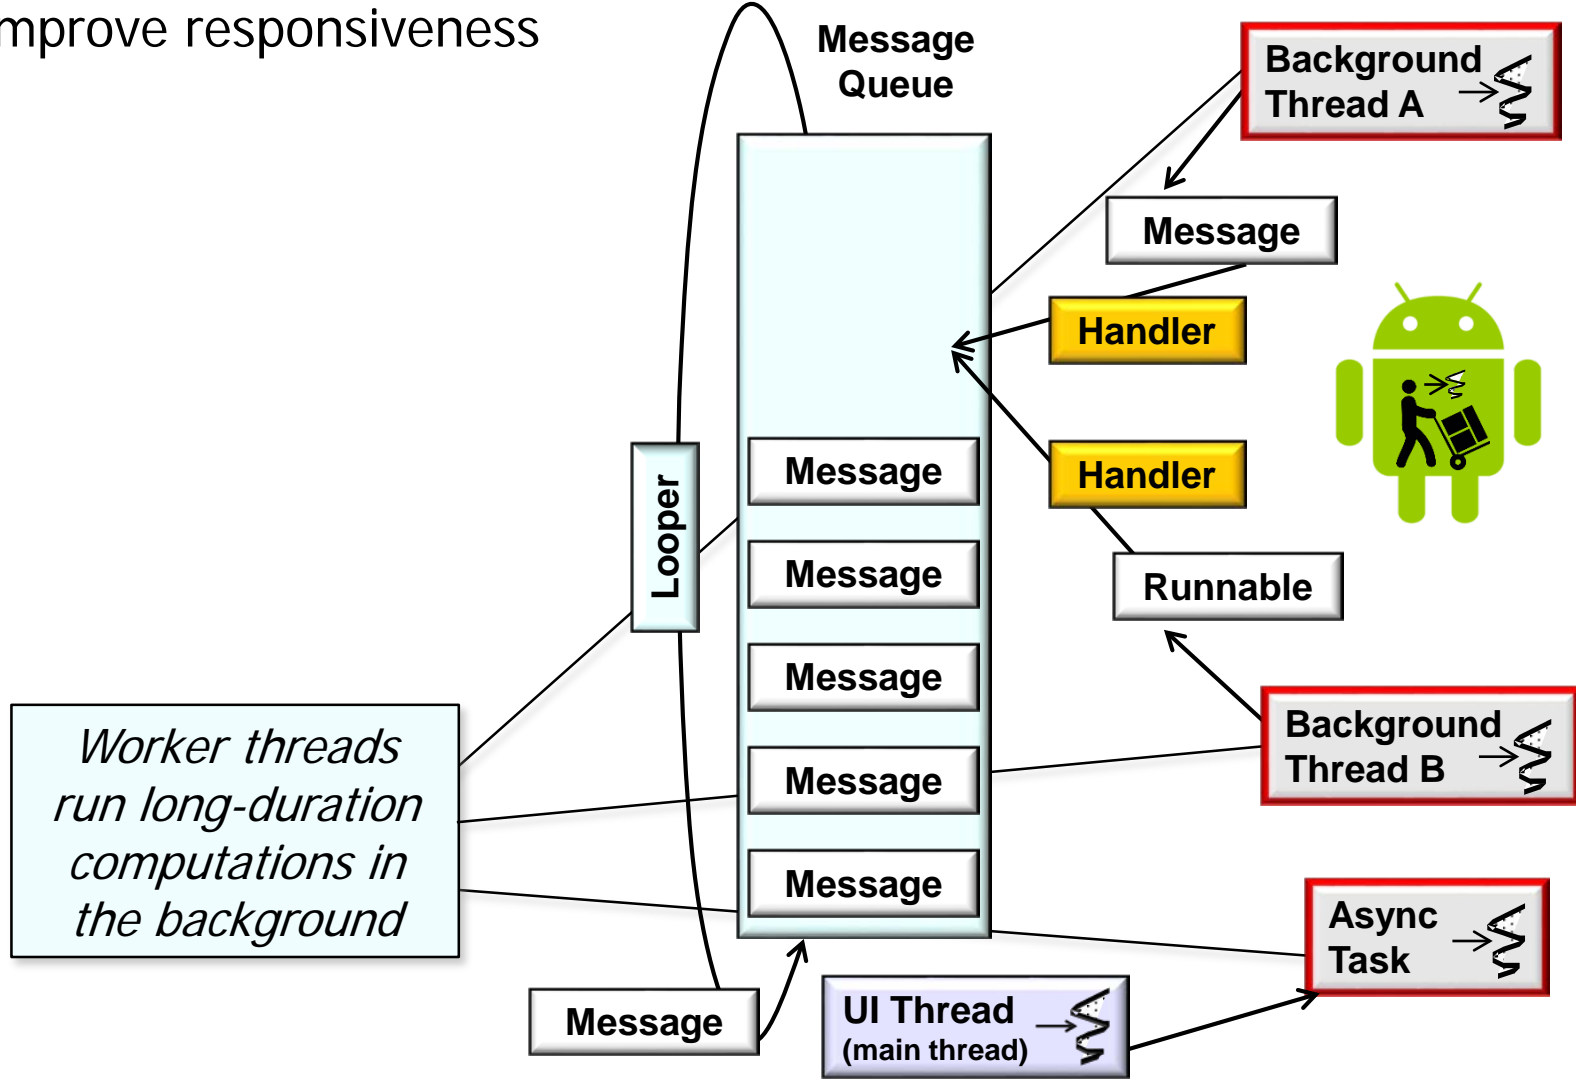
Using Concurrency to Improve Responsiveness

- Android concurrency frameworks define features & idioms that can help ensure & improve responsiveness



Using Concurrency to Improve Responsiveness

- Android concurrency frameworks define features & idioms that can help ensure & improve responsiveness



Worker threads can block waiting for I/O or computations to complete

End of Benefits of Concurrency in Java & Android (Part 2)